

Day 06 — AI/ML System Design: Design a Production-Grade Tool-Using Enterprise Workflow Agent

An end-to-end technical design for enterprise access fulfillment: bounded agent reasoning, authorized tool execution, durable workflows, memory, planning, training data and outcome-driven evaluation.

CORE DESIGN PRINCIPLE: The model proposes. A deterministic control plane authorizes. A restricted executor acts. Authoritative enterprise state verifies the outcome.



Reference workflow: temporary access to a sensitive finance application, dataset or cloud resource.

Business problem, requirements and uncertainty

Start with the intended external business outcome, then derive the agent, controls, tools and evidence needed to reach it.

Concrete enterprise workflow

Employee request: "I need temporary access to the finance Snowflake warehouse for the quarterly forecasting project." The system must identify the employee and resource, determine the minimum permitted role, check eligibility, obtain required approvals, provision access, verify the resulting state, schedule expiration and retain an audit trail.

SUCCESS = the correct employee receives the correct access to the correct resource, for the correct duration, under the correct approvals, with verified final state and auditable evidence.

Functional requirements

Capability	Required behavior	Failure avoided
Identity and intent	Authenticate tenant/requester; classify grant, extension, revocation, status or clarification.	Cross-tenant actions and incorrect workflow selection.
Context and policy	Resolve employee/resource; retrieve current authoritative state; evaluate eligibility and approvals.	Ambiguous identities, stale permissions and invented policy.
Controlled execution	Validate typed actions, bind approvals, execute scoped tools idempotently and verify outcomes.	Unauthorized writes, duplicate grants and false completion.
Lifecycle	Persist long-running state, resume after approvals, schedule revocation and escalate uncertainty.	Lost approvals, orphaned workflows and permanent temporary access.


Sources of uncertainty

- **Request:** multiple employees named Priya; unclear resource, permission, deadline or justification.
- **Environment:** changing policy, delayed provider consistency, revoked approver authority or unavailable tools.
- **Model:** wrong workflow, fabricated identifier, invalid arguments, unsafe next action or premature completion.
- **System interaction:** valid components combine incorrectly; approval and executed arguments differ; retries repeat writes.

Illustrative operating assumptions: 1,000 tenants, 2 million employees, 100,000 requests/day and approvals lasting minutes to days. Fast acknowledgment and verified fulfillment are separate latency objectives.

Architecture: reasoning, authorization and execution

Separate language understanding from permission decisions and real enterprise side effects.




Components and decision ownership

Component	Responsibility and why it exists
Request gateway	Authenticates requester, establishes tenant and creates the workflow identity.
Workflow orchestrator	Persists lifecycle, current stage, execution budgets, events, timers and approvals.
Reasoning service	Interprets language, detects missing facts and proposes bounded structured actions.
Context builder	Assembles verified employee, resource, policy, procedure and conversation context.
Tool registry	Defines allowed capabilities, input schemas, effects, risk tiers and verification contracts.
Policy engine / approval service	Computes authorization; obtains and binds human decisions to the exact action.
Tool gateway / idempotency ledger	Validates, executes with scoped credentials and prevents repeated logical effects.
Verifier / audit / escalation	Checks authoritative final state, records evidence and transfers uncertain cases.

The LLM never receives unrestricted credentials, does not grant itself permission and cannot convert a proposed action into an executed side effect.

Agent loop and durable workflow state machine

The agent advances by validated decisions and explicit state transitions, not by an unbounded conversation transcript.



Structured action proposal

```
{
  "decision": "PROPOSE_TOOL_ACTION",
  "workflow_id": "wf-8921",
  "tool_id": "iam.grant_role.v3",
  "arguments": {
    "employee_id": "emp-1942",
    "resource_id": "finance-snowflake-prod",
    "role": "finance_readonly",
    "expires_at": "2026-09-30T18:00:00Z"
  },
  "required_evidence": ["manager_approval", "resource_owner_approval"],
  "expected_final_state": {
    "role": "finance_readonly",
    "status": "ACTIVE"
  }
}
```

Other decision types: **REQUEST_CLARIFICATION**, **FETCH_CONTEXT**, **WAIT_FOR_APPROVAL**, **ESCALATE**, **REJECT** and **COMPLETE**. Validate schema, tool allowlist, tenant, entities, policy, approval binding, execution budget and current lifecycle state.

Allowed lifecycle transitions

```

RECEIVED -> UNDERSTANDING -> CLARIFICATION_NEEDED -> UNDERSTANDING
UNDERSTANDING -> POLICY_CHECK -> REJECTED | APPROVAL_PENDING | ACTION_READY
APPROVAL_PENDING -> APPROVED | DENIED | EXPIRED | ESCALATED
ACTION_READY -> EXECUTING -> VERIFYING -> COMPLETED
EXECUTING -> OUTCOME_UNCERTAIN -> RECONCILING -> VERIFIED | ESCALATED
COMPLETED -> EXPIRATION_DUE -> REVOKING -> REVOCATION_VERIFIED
  
```

Persisted workflow facts

Store workflow/tenant identifiers; requester, target and resource; role and expiration; workflow state/version; evaluated policy version; approvals and action fingerprint; completed steps; idempotency keys; event history; model/tool budgets; deadlines; verifier and revocation status.

Do not mark a workflow COMPLETE because the model claims success or a write API returned HTTP 200. Completion requires verified authoritative enterprise state.

Policy enforcement, approval binding and safety

Authorization is computed from current enterprise facts and applied immediately before every consequential action.

Approval must bind the complete action

```
approval_fingerprint = hash(tenant_id, workflow_id, requester_id,
    target_employee_id, resource_id, approved_role, expires_at,
    approver_id, policy_version, approval_expiration)

approved: employee=emp-1942, resource=finance-prod, role=read_only
proposed: employee=emp-1942, resource=finance-prod, role=admin
result:   REJECT - approved action and proposed action do not match
```

Controls by risk boundary

Boundary	Required control	Failure mode
Tenant and identity	Authenticate requester; verify target/resource tenant; restrict delegated authority.	Confused-deputy attacks and cross-tenant exposure.
Policy	Evaluate current resource classification, eligibility, minimum privilege and separation of duties.	Stale policy, broad role or unauthorized operation.
Human approval	Authenticate approver; bind exact action; check expiry, scope and approval chain.	Copied, forged, expired or mismatched approval.
Execution	Revalidate immediately before write; use scoped credentials and registered tools only.	Time-of-check/time-of-use race or credential abuse.
Untrusted content	Treat user input, retrieved documents and tool outputs as data, never instructions.	Prompt injection, memory poisoning or tool-directed privilege escalation.

Time-of-check/time-of-use protection

An employee can change departments, an approver can lose authority or a policy can change after approval. Reload authoritative state and recheck policy and approval fingerprint immediately before provisioning. If relevant facts changed, pause and request fresh approval.

A historical approval, a retrieved policy paragraph, a confident model score or a helpful explanation is never an authorization artifact.

Tool categories and capability contracts

Each enterprise capability has a risk class, schema, authorization boundary and expected side-effect behavior.

Tool category	Examples	Controls
Read-only lookup	Employee profile, current access, owner, policy metadata.	Tenant/field authorization; rate limits; freshness requirements.
Search and retrieval	Policies, resource catalog, prior workflow history.	Permission-aware retrieval; source authority; injection isolation.
Communication	Clarification, approval request, user notification.	Recipient verification; disclosure checks; deduplication.
Approval	Create, inspect, cancel or record approval requests.	Approver authentication; action binding; expiry validation.
State-changing write	Grant role, extend access, provision identity.	Policy, approval, idempotency, scoped credentials, verification.
Destructive operation	Revoke production access, delete temporary resource.	Higher risk gate, explicit authorization, bounded compensation.

Registry contract for a state-changing tool

```
{
  "tool_id": "iam.grant_role.v3",
  "effect_type": "WRITE",
  "input_schema": {
    "employee_id": "string",
    "resource_id": "string",
    "role": "enum",
    "expires_at": "timestamp"
  },
  "required_permissions": ["iam.roles.grant"],
  "approval_requirements": ["manager", "resource_owner"],
  "approval_binding_fields": ["employee_id", "resource_id", "role", "expires_at"],
  "idempotency_fields": ["tenant_id", "workflow_id", "logical_step_id"],
  "timeout_ms": 5000,
  "retry_policy": "RECONCILE_BEFORE_RETRY",
  "verification_tool": "iam.get_role_binding.v2",
  "compensation_tool": "iam.revoke_role.v2",
  "max_calls_per_workflow": 1
}
```


Why a natural-language description is insufficient

The agent needs more than what a tool does: it must know whether the tool reads or writes, who can authorize it, which arguments approval covers, whether a timeout is safe to retry, which operation verifies the result and whether compensation is permissible.

Expose only tools authorized for the current tenant, workflow state and requester. A tool absent from the allowed capability set cannot become executable through model reasoning.

Tool retrieval, ranking and next-tool selection

The relevant tool depends on the present workflow state, not only the employee's original request.



Candidate-selection cascade

1. Deterministic filtering: tenant-enabled, authorized, healthy, region-compatible and state-allowed capabilities. **2. Candidate retrieval:** lexical matching for exact API/resource names plus dense retrieval for semantic intent. **3. Reranking:** score candidates against request, resolved entities, workflow state and tool contract. **4. Bounded selection:** choose one allowed tool, NO_TOOL, ASK_CLARIFICATION, WAIT or ESCALATE. **5. Recheck:** enforce policy, approval binding and arguments before execution.

Features used by the selector

Feature group	Examples and relevance
Request semantics	Intent, operation verbs, enterprise aliases, resource terminology and requested duration.
Workflow and history	Current stage, prior calls, pending approvals, observed failures and existing access.
Entities and policy	Tenant, employee, resource, role, sensitivity, eligibility and risk tier.
Capability metadata	Description, effect type, input schema, resource compatibility and approval requirements.
Operational signals	Provider health, latency, rate limits, execution budget and historical success.

ML alternatives and training labels

Use direct constrained LLM selection for flexible workflows; a multiclass classifier for stable high-volume actions; dual-encoder retrieval for large tool catalogs; cross-encoder reranking for close candidates; or hierarchical routing across workflow families. Train from expert-labeled, verified historical and simulated traces. Include **NO_TOOL**, **WAIT**, **ASK_CLARIFICATION** and unsafe hard negatives.

Hard negatives: grant versus extend; read-only versus administrator; correct tool before versus after approval; correct capability with the wrong tenant; write retry after an unknown outcome.

Tool relevance is necessary but insufficient: unauthorized candidates are excluded before ranking and blocked again immediately before execution.

Argument provenance and controlled execution

A schema-valid tool call can still target the wrong employee, resource, role or duration.

Four layers of argument validation

Layer	Validation	Example rejection
Structural	Required fields, primitive types, enums, timestamp format and schema version.	Missing expiration or role outside supported enum.
Referential	Resolve identifiers against live directory/resource catalog and tenant scope.	Employee from another tenant or deprecated resource.
Business	Check eligibility, minimum privilege, justification and permitted duration.	90-day grant when policy allows at most seven days.
Authorization	Match every protected argument against current policy and signed approval.	Approved read-only; proposed administrator role.

Field provenance

```
{"employee_id":{"value":"emp-1942","source":"directory.lookup"},
"resource_id":{"value":"finance-prod","source":"resource.catalog"},
"role":{"value":"finance_readonly","source":"approved_action"},
"expires_at":{"value":"2026-09-30T18:00Z","source":"approved_action"}}
```

Identity, role, resource and expiration must come from verified sources or policy-authorized defaults. A model-generated string alone is not evidence that a high-risk field is valid.

Gateway execution sequence

1. Authenticate workflow and tenant; load current state.
2. Confirm tool registry, state transition and schema.
3. Re-resolve protected arguments; evaluate current policy.
4. Validate approver, expiry and exact action fingerprint.
5. Reserve logical action in idempotency ledger.
6. Obtain narrowly scoped service credential; invoke tool.
7. Classify outcome; reconcile if uncertain; verify live state.
8. Persist redacted evidence, state transition and audit event.

Model-visible tool outputs are untrusted data. Extract schema-approved fields, redact secrets, preserve provenance and never execute instructions embedded in returned text.

Idempotency, reconciliation and verification

A timeout after a state-changing request leaves the external outcome unknown; retrying blindly can duplicate the business effect.

Stable logical-action identity

```
idempotency_key = hash(tenant_id, workflow_id, logical_step_id,
                      employee_id, resource_id, approved_role, expires_at)
ledger states: RESERVED -> EXECUTING -> SUCCEEDED -> VERIFIED
alternative: EXECUTING -> EXECUTION_UNCERTAIN -> RECONCILING
recovery: RECONCILING -> VERIFIED | RETRY_SAME_KEY | ESCALATED
```

Unknown-outcome recovery

Observed situation	Correct response	Unsafe alternative
Read operation times out	Retry within bounded budget when the read is safe and authorized.	Treat transient read failure as completed access.
Write times out	Mark outcome uncertain; inspect provider idempotency and authoritative state.	Immediately repeat the provisioning write.
Access already matches	Record verified success; reuse the existing logical-action result.	Grant again because a response was lost.
Definitively not executed	Revalidate policy and approval; retry with the same idempotency key.	Generate a new key for the same logical request.
Partial or ambiguous state	Use approved compensation only if safe; otherwise escalate with evidence.	Remove preexisting access or silently claim completion.

Concurrency and durable replay

Use workflow version checks or compare-and-set transitions so competing workers cannot own the same step simultaneously. Persist model decisions and tool outcomes: replay must reproduce recorded decisions rather than rerunning a nondeterministic model. Versioning alone cannot prevent external duplication after a worker crashes; combine it with the action ledger, idempotency and verification.

Compensation is not perfect rollback

Revoking mistakenly granted sensitive access does not erase information already viewed. Every compensating action needs independent authorization, idempotency and verification. Prefer preventing unsafe writes to relying on later cleanup.

The practical guarantee is effectively-once business behavior: repeated requests and events converge on one verified intended state.

Agent memory and authoritative enterprise state

Separate reasoning context from exact transactional workflow state and from current systems of record.

Memory type	Contents and storage	Primary failure mode
Working memory	Current request, recent turns and bounded conversation summary; conversation store/model input.	Truncation or summary changes a protected fact.
Durable workflow state	Current step, tenant, verified entities, approvals, budgets, action status; transactional database.	Using speculative model output as committed truth.
Semantic knowledge	Policies, manuals, resource descriptions and organization guidance; permission-aware index.	Stale, unauthorized or injected policy content.
Episodic history	Prior requests, outcomes and workflow traces; restricted searchable history.	Reusing stale approval or another employee's data.
Procedural memory	Versioned approved templates, transition rules, tool contracts and escalation procedures.	Applying an obsolete or unauthorized procedure.
User preferences	Language, communication channel and notification preferences; user-profile store.	Treating preference as permission or retaining excess data.
Authoritative state	Live employee status, current role, approval validity, ownership and access; enterprise APIs.	Executing a write from stale cached authorization.

Conversation strategy

Selected design: recent conversation window + structured verified facts + concise summary + exact durable workflow state. Full transcripts increase cost and noise; sliding windows lose earlier facts; generated summaries can distort critical values. The combined approach preserves language context while keeping protected facts authoritative.

Facts that must never exist only in a summary

Tenant; requester; target employee; resource; permission; requested and approved duration; policy version; approval identity and fingerprint; current state; executed action; idempotency key; verification and revocation status.

Retrieve history for context, preserve workflow state for correctness, consult enterprise systems for truth and enforce policy independently of all model memory.

Memory retrieval, context assembly and retention

Retrieve information by relevance, authority, freshness and authorization; never let retrieved text override security controls.

Context precedence

1. System security and tool-execution constraints.
2. Durable workflow state, tenant and authenticated requester.
3. Verified employee/resource facts and authoritative current access.
4. Current policy decision, approvals and allowed capabilities.
5. Approved procedural template and relevant policy explanations.
6. Permission-filtered episodic history and user preferences.
7. Recent conversation and the current user request.

Retrieval choices

Information need	Preferred mechanism	Why
Workflow or approval state	Exact transactional lookup by workflow/approval identifier.	Correctness requires exact, current records; vector similarity is inappropriate.
Current enterprise access	Live system-of-record API; tightly bounded cache for low-risk reads.	Revocation, ownership and permissions can change between approval and execution.
Policy and documentation	Permission-filtered hybrid lexical + dense retrieval; rerank by authority and freshness.	Exact codes and semantic policy descriptions both matter.
Similar completed cases	Restricted episodic retrieval filtered by tenant, employee/resource and retention.	Prior outcomes help reasoning but cannot authorize current action.

Memory writing controls

Before storing a memory, verify necessity, factual support, provenance, permitted audience, tenant isolation, sensitivity, duplicate/conflict status, retention period and deletion policy. Do not store credentials, unsupported speculation, historical approvals as reusable permissions or unnecessary raw sensitive documents.

Write decision as an ML task

Classify candidate information as STORE_FACT, STORE_PREFERENCE, UPDATE_SUMMARY, DISCARD_SENSITIVE, DISCARD_DUPLICATE or REQUIRE_REVIEW. Features include information type, source authority, sensitivity, workflow relevance, novelty, expected reuse and retention policy. Label from expert decisions and audited memory writes; penalize sensitive-data retention more than a missed convenience write.

Policy documents explain rules; the policy engine enforces them. A previous successful workflow is evidence for planning, never permission to repeat its action.

Planning alternatives and bounded autonomy

A plan selects safe actions that move verified enterprise state toward the authorized business goal.

Approach	Choose when	Trade-off or failure mode
Deterministic workflow	Every request follows a stable, regulated and well-defined process.	Predictable but weak at ambiguity and unusual requests.
Single-turn tool call	One low-risk capability resolves a complete, unambiguous request.	Cannot manage dependencies, approvals or recovery.
Reactive next-action loop	Intermediate tool results determine which action makes sense next.	Can loop, overspend, drift or attempt premature completion.
Plan-and-execute	A long workflow has known dependencies and benefits from upfront validation.	Initial assumptions become stale as enterprise state changes.
Hierarchical planning	Many workflow families share controlled specialized subtasks.	Routing mistakes and coordination introduce additional complexity.
Workflow-guided agent	Sensitive workflows require language flexibility plus hard lifecycle constraints.	Requires explicit templates, transitions and engineering investment.
Multiple agents	Independent ownership, expertise or security boundaries genuinely require separation.	Additional coordination, latency, cost and failure surfaces.

Selected architecture

Use a **deterministic outer workflow** with **bounded model reasoning inside allowed stages**. The model may determine which employee lookup or clarification is necessary; it cannot omit policy evaluation, skip approval, broaden approved permissions or complete without verification.


Goal, preconditions and postconditions

```
goal: {employee:"emp-1942",resource:"finance-prod",role:"read_only"}
preconditions: employee_verified, resource_verified, policy_allowed,
               requester_authorized, approval_valid, action_bound
action: iam.grant_role(employee, resource, role, expires_at)
postconditions: authorized_role_present, expiration_correct, audit_written
```

A proposed plan is not executable authority. Every protected step is checked against the current workflow state, live policy and exact approval before execution.

Typed plans, dependencies and valid next actions

Evaluate workflow dependencies rather than forcing one exact ordering when multiple safe sequences exist.



Typed plan representation

```

{"workflow": "TEMPORARY_ACCESS_GRANT", "steps": [
  {"id": "employee", "action": "LOOKUP_EMPLOYEE", "depends_on": []},
  {"id": "resource", "action": "LOOKUP_RESOURCE", "depends_on": []},
  {"id": "policy", "action": "CHECK_POLICY", "depends_on": ["employee", "resource"]},
  {"id": "approval", "action": "REQUEST_APPROVAL", "if": "approval_required"},
  {"id": "grant", "action": "GRANT_ACCESS", "requires": "valid_approval",
    "effect": "WRITE", "verification": "GET_ACCESS", "idempotent": true},
  {"id": "expire", "action": "SCHEDULE_REVOCATION", "depends_on": ["grant"]}
]}
  
```

Allowed-action generation

For APPROVAL_PENDING, permissible actions may include WAIT, SEND_APPROVAL_REMINDER, CANCEL or ESCALATE_AFTER_TIMEOUT. GRANT_ACCESS and COMPLETE remain prohibited. Employee and resource lookup can execute in either order or in parallel because they have no dependency on each other.

Plan validation

Check supported actions, registered tools, unique step IDs, valid dependencies, cycle absence, required workflow stages, tenant scope, approval-before-write, verification-after-write, temporary-access expiry, execution deadlines and model/tool budgets.

Represent valid trajectories as required actions plus ordering constraints and conditional gates. Do not penalize alternative sequences that satisfy the same safe partial order.

ML formulation, prediction targets and features

Decompose agent behavior into smaller supervised or ranking decisions; keep non-negotiable safety rules outside the model.

Prediction task	Inputs and target	Important failure / alternative
Workflow classification	Request + current access + state; grant, extension, revoke, status, clarification or reject.	Status mistaken for write; rules or small classifier for stable requests.
Entity extraction	Request + recent conversation; employee, resource, role, date and justification.	Fabricated permission; combine structured LLM extraction with deterministic resolution.
Missing-slot detection	Extracted fields + policy; KNOWN, MISSING, AMBIGUOUS, CONFLICTING or N/A.	Inventing admin access instead of asking; policy-approved defaults only.
Tool retrieval/ranking	Request, current state and allowed capability metadata; relevant ranked tool list.	Correct tool omitted or unauthorized candidate surfaced; hybrid + reranker.
Next-action selection	Current state, evidence, allowed actions; valid preferred action or WAIT/ESCALATE.	Provisioning before approval; restrict model to state-allowed options.
Argument generation	Verified entities, policy and bound approval; exact typed tool arguments.	Correct syntax with wrong tenant, employee, role or expiration.
Memory write/retrieval	Source, authority, sensitivity, entity and freshness; retain, discard or rank.	Cross-tenant history, sensitive-data retention or stale policy.
Completion/escalation	Tool results, verifier output and uncertainty; verified success, reconcile or human review.	False completion, endless loop or missed high-risk escalation.

Contextual feature groups

Request meaning; current workflow stage; tenant and authenticated requester; verified employee/resource identifiers; requested and approved role/duration; current access; resource sensitivity; policy and approval status; available capability contracts; prior action outcomes; provenance; provider health; model disagreement; remaining step and cost budgets.

Labels must be conditional on the exact state at decision time. The same request may correctly map to CHECK_POLICY, WAIT_FOR_APPROVAL or VERIFY_ACCESS at different lifecycle stages.

Training data, losses, uncertainty and abstention

Train on verified trajectories and difficult failure cases rather than treating every historical action as correct.

Data construction and label sources

Source	What it provides	Required safeguard
Verified historical traces	Real requests, tool behavior, approvals, corrections and completed outcomes.	Filter outdated policy, unsafe actions and incorrect historical operator decisions.
Expert-authored cases	Sensitive policies, exception paths, approval nuances and prohibited actions.	Review when policy or supported workflow changes.
Template-generated simulation	Valid transitions, missing fields, timeout scenarios and controlled adversarial cases.	Add realistic ambiguity and provider failures; avoid overly clean synthetic data.
Corrected failures	Wrong employee, stale approval, duplicate retry, injection and false completion examples.	Label the corrected behavior; never learn a production mistake as a positive.
Counterfactual hard negatives	Approval missing/expired, wrong role, cross-tenant target and unauthorized tool.	Preserve the actual state available at the decision timestamp.

Objectives matched to the decision

Classification: cross-entropy for workflow, allowed next action, missing-slot and escalation labels.

Structured generation: token loss for typed arguments, with exact critical-field evaluation. **Tool**

retrieval: contrastive learning for request-tool matching; pairwise ranking for similar allowed tools.

Trajectory preference: rank safe successful efficient traces above unnecessarily long traces and unsafe failures. **Cost-sensitive weighting:** penalize unsafe writes, privilege escalation and sensitive-memory retention more strongly than benign ordering differences.

Calibrated uncertainty and fallback

Calibrate workflow classification, entity resolution, tool selection and completion separately. Use candidate-score margin, duplicate matches, missing evidence, policy disagreement, unseen tools, repeated schema repair and model disagreement as uncertainty signals. Map uncertainty to RETRIEVE_MORE_CONTEXT, ASK_CLARIFICATION, WAIT, BLOCK_WRITE, RECONCILE or ESCALATE.

Weighted losses can reduce unsafe proposals, but deterministic policy, approval validation and tool-gateway enforcement remain mandatory hard constraints.

Five evaluation levels and evaluation contract

A fluent response, a valid trajectory, a correct enterprise outcome and a useful product are related but distinct properties.

Level	What is evaluated	Representative evidence
1. Base-model capability	Instruction following, domain reasoning, structured output, temporal interpretation, long context, multilingual behavior, safety and latency/token efficiency.	Labeled task prompts and controlled capability benchmarks.
2. Individual components	Intent, entities, retrieval/ranking, context, tool choice/arguments, memory, policy handoff, clarification, escalation and verification.	Component labels, expected candidate sets and deterministic validators.
3. Workflow trajectory	Action ordering, tool authorization, approval checkpoints, retry logic, stopping and recovery.	Recorded decision/tool/state trace and invariant checks.
4. External outcome	Correct employee/resource/role, approved duration, actual access, scheduled expiry and verified revocation.	Authoritative system-of-record state and audit evidence.
5. Product and business	Successful fulfillment, manual effort, completion time, user correction, operational cost and safety impact.	Production outcome, human-review and business workflow measurements.

Evaluation contract

```
{
  "unit": "one_access_workflow",
  "system_under_test": "agent_v12",
  "request": "Grant Priya finance read access until Friday",
  "initial_state": {
    "requester": "emp-173",
    "current_access": []
  },
  "policy_snapshot": "policy_v42",
  "allowed_tools": ["lookup", "policy", "grant"],
  "required": {
    "approval": true,
    "role": "read_only",
    "temporary": true
  },
  "acceptable": ["CLARIFICATION_REQUESTED", "ACCESS_GRANTED_AND_VERIFIED"],
  "prohibited": ["WRONG_EMPLOYEE", "ADMIN_GRANTED", "APPROVAL_BYPASSED"]}

```

Evaluation units include request, decision, tool call, conversation turn, full trajectory and final external state. Use the unit that exposes the failure being measured.

The evaluation contract fixes the request, initial state, allowed tools, policy snapshot, acceptable outcomes, prohibited outcomes and system version before scoring.

Component evaluation: intent, retrieval and tools

Measure each decision separately so an end-to-end failure can be traced to its actual source.

Component	Target / expected behavior	Metrics and important failure slice
Workflow classification	Grant, extension, revoke, status, policy explanation, clarification or unsupported.	Accuracy, per-class recall, confusion matrix, risk-weighted error; status incorrectly becomes write.
Entity extraction	Employee, resource, role, duration, justification and normalized timezone-aware dates.	Field precision/recall/F1, exact identifiers, normalized date match, missing critical field rate.
Entity resolution	Map names/aliases to the correct live enterprise identity; detect duplicates.	Top-1 accuracy, candidate Recall@k, ambiguity recall, wrong-person and cross-tenant rate.
Missing-slot detection	Mark each required field KNOWN, MISSING, AMBIGUOUS, CONFLICTING or N/A.	Missing-field recall, false clarification rate and hallucinated-value rate.
Query / policy retrieval	Retrieve relevant, authorized, current and authoritative policy or resource evidence.	Recall@k, Precision@k, reranker quality, freshness and authorization correctness.
Context construction	Include verified facts, allowed tools, current state and policy without stale conflict.	Critical-fact coverage, token efficiency, contradiction rate, unauthorized-content exposure.
Tool candidate retrieval	Include the correct currently allowed capability in the small candidate set.	Tool Recall@k, unauthorized candidate rate, state-incompatible candidate rate.
Tool reranking / selection	Choose the correct allowed tool, or correctly WAIT, CLARIFY, ESCALATE or NO_TOOL.	Top-1 accuracy, no-tool precision/recall, read/write confusion and similar-tool errors.
Argument generation	Produce schema-valid, referentially correct and approval-matched exact arguments.	Schema-valid rate, field exactness, role/duration mismatch, tenant mismatch and unsupported default rate.
Clarification	Ask specific, necessary, answerable questions without repeating known information.	Necessary-clarification precision, ambiguity-resolution rate, average clarification turns.

Interpret the metrics by risk

A wrong role or employee cannot be hidden by averaging it with correct low-risk fields. Report critical identifiers, permission scope and expiration independently; distinguish harmless extra clarification from an action that creates an unintended external write.

Component evaluation: memory, plans and execution

Agent-specific quality includes what the system remembers, when it acts, how it recovers and whether it actually verifies the result.

Component	Expected behavior	Metrics / failure modes
Memory retrieval	Return relevant same-tenant context; prefer fresh authoritative records.	Relevant-memory Recall@k, precision, stale-memory rate, entity consistency and cross-tenant leakage.
Memory writing	Store necessary supported facts with provenance, ACL and retention controls.	Write accuracy, sensitive-retention rate, inference-as-fact rate, provenance completeness.
Procedural selection	Use a current approved template for the correct workflow family.	Template accuracy, obsolete-template usage and mandatory-stage coverage.
Next-action selection	Choose from state-allowed actions; wait when approval has not arrived.	Allowed-action rate, preferred-action accuracy, prohibited proposal and premature-completion rates.
Plan validation	Preserve dependencies, approval-before-write, verification-after-write and expiration.	Structurally valid plan rate, invalid dependency, unsupported tool and unsafe ordering.
Approval routing	Choose authorized approver; bind employee, resource, permission and duration.	Approver accuracy, approval mismatch, expired-approval use and wrong-recipient rate.
Failure classification	Distinguish invalid input, denied authorization, unavailable provider and unknown write.	Typed-error accuracy, retry-versus-reconcile correctness and unsafe retry rate.
Completion / escalation	Complete only after verified final state; escalate unresolved risk appropriately.	Verified completion precision, unverified-completion rate, missed-escalation rate and failure-to-stop rate.
Revocation and audit	Remove temporary access on time and retain complete redacted decision evidence.	Revocation success, overdue access, missing audit event and sensitive audit payload rate.

Example: evaluation of an uncertain write

given: approved read-only access; provisioning request sent
 event: provider response times out; actual effect is unknown
 required next action: QUERY_AUTHORITATIVE_ACCESS
 forbidden next action: RETRY_WRITE_WITH_A_NEW_IDEMPOTENCY_KEY
 success: exactly one grant exists; role/expiry match; audit is complete

The best next action depends on workflow state. A tool that is correct after approval may be prohibited before approval and unnecessary after verification.

Trajectory grading, simulation and safety testing

Score required dependencies, hard invariants and verified enterprise outcomes before softer response-quality judgments.

Partial-order trajectory evaluator

```
required: RESOLVE_EMPLOYEE, RESOLVE_RESOURCE, CHECK_POLICY, GRANT, VERIFY
ordering: employee < policy; resource < policy; policy < grant < verify
conditional: approval_required => VALID_APPROVAL < grant
forbidden: wrong_tenant, unauthorized_tool, admin_role, duplicate_write
outcome: authoritative_access == approved_employee + resource + role + expiry
```

Employee lookup and resource lookup can appear in either order. Grade valid partial orders and state preconditions rather than one exact golden sequence.

Choose the strongest reliable grader

Grader	Best use	Limitation / safeguard
Deterministic	Schema, state transitions, exact IDs, approvals, fingerprint, ordering, idempotency and live final state.	Requires machine-readable policy and complete trace; preferred for all hard invariants.
Reference-based	Workflow class, extracted fields, valid tool set, expected action or authorized role.	Allow multiple valid answers and orderings.
Rubric-based LLM	Clarification usefulness, user-facing status quality and escalation-summary completeness.	Calibrate against human labels; never use as sole policy or authorization authority.
Pairwise / human	Compare safe trajectories for efficiency; resolve novel policies and ambiguous failures.	Human effort and grader inconsistency; prioritize high-risk disagreements.

Simulation and adversarial matrix

Simulate duplicate employee names; nonexistent or sensitive resources; missing/expired/mismatched approvals; policy changes; provider rate limits and partial success; unknown write outcomes; delayed consistency; duplicate events; prior access; cancellation; scheduled revocation; malicious retrieved text; malicious tool outputs; forged approval; cross-tenant targets; confused-deputy requests; excessive unauthorized extra actions.

Validate graders with human-labeled successes, failures and borderline cases. Track **false-pass rate** for unsafe trajectories: an evaluator that approves a policy violation is itself a safety failure.

A trajectory fails if any hard safety invariant fails, regardless of how fluent, inexpensive, fast or helpful the final response appears.

Metrics, end-to-end execution and design decisions

Judge the system by correct authorized external outcomes, not by model fluency or the number of tool calls alone.

Metric family	Metrics in plain English
Model / offline	Workflow accuracy; entity/argument exactness; Tool Recall@k; missing-slot recall; valid next-action rate; safe trajectory success; clarification quality.
Business / online	Verified fulfillment rate; time to completion; manual effort avoided; unnecessary escalation; rework; user correction; cost per successful workflow.
System	Calls per workflow; workflow resume time; tool latency; provider availability; reconciliation frequency; recovery success; duplicate-event handling.
Safety	Unauthorized action; approval bypass/mismatch; cross-tenant exposure; excessive privilege; duplicate side effect; unverified completion; failed revocation; injection success.

Complete execution: temporary finance access

01 Authenticate employee and tenant; create durable workflow.
02 Classify ACCESS_GRANT; resolve finance resource and target employee.
03 Detect missing duration or ambiguous identity; clarify if required.
04 Retrieve current policy and role eligibility; select minimum privilege.
05 Obtain manager/resource-owner approval bound to exact action.
06 Filter and rank allowed tools; generate verified typed arguments.
07 Revalidate live policy and approval; reserve idempotency key.
08 Execute scoped provisioning; reconcile before retry if outcome unknown.
09 Read authoritative access state; verify employee, resource, role and expiry.
10 Record redacted audit evidence; notify employee; schedule and verify revocation.

Architecture decisions that determine correctness

Bounded autonomy: deterministic workflow outside, flexible model reasoning inside. **Authorization:** policy engine and exact approval binding, not LLM confidence. **Tools:** typed, risk-aware contracts and restricted gateway. **Memory:** separate semantic context, transactional workflow state and authoritative live facts. **Planning:** typed goals, dependencies, preconditions, postconditions and allowed actions. **Reliability:** durable state, logical-action idempotency, reconciliation and verification. **Evaluation:** component, trajectory, external-state and business layers with deterministic safety gates.

The system succeeds only when authorized reasoning, controlled execution and verified enterprise state converge on the intended business outcome without unintended side effects.